

Assignment 4 Report

By Aniruddha Bala
Roll No: 15655

December 1, 2018

1 Parallelization of KMeans Clustering using OpenMP

1.1 Methodology

1.1.1 Visualizing the data


At first the data was read from input file and the data points were plotted in python using matplotlib. The data was found to have a good natural structure arranged in 20 blocks. The plot obtained is shown below.

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
In [98]: data = pd.read_csv('../openmp/data_500k.csv', header=None)
```

```
In [99]: plt.figure(figsize = (20,20))
plt.scatter(data[0],data[1])
plt.xlabel('x')
plt.ylabel('y')
```

```
Out[99]: Text(0,0.5,'y')
```



1.1.2 Sequential Kmeans

Then I implemented the sequential program for KMeans Algorithm. The algorithm takes n points as input and forms K clusters out of it. In this case we had 500k points and we had to group them into 20 clusters. The KMeans algorithm works as follows:

1. Initialize random centroids for each of the K clusters
2. Repeat until convergence do:
3. Cluster assignment step: For each of the n points assign them to a cluster(i) whose centroid is closest to the point.
4. Move centroid step: For each of the clusters update the cluster centroid to the average of points assigned to that cluster.

1.1.3 Parallelizing with OpenMP

In each loop the major computation in the Kmeans algorithm is done in the part where we assign cluster centroids to each of the n points. Therefore parallelizing that loop makes sense. The loop has been parallelized using the openmp for directive. The parallelization has been tried for 4, 8, and 16 number of threads and using static and dynamic scheduling for assigning iterations to threads.

1.2 Experimental setup

Input data: The input data file data_500k.csv contains the (x,y) coordinates of 500000 points.

Output : The program outputs 20 lines each line containing the cluster index of the cluster formed, the number of points in the cluster, centroid coordinates of the cluster. At the end the execution time of the program is also printed.

The following steps were followed while conducting the experiment:

1. Run the sequential program, note the result and the execution time.
2. Run the parallelized openmp version with all possible combinations of number of threads (4, 8, 16) and schedule clause (dynamic and static)
3. Note the execution time in each case.
4. To maintain uniformity both the sequential and parallelized versions of the code have been timed with openmp's high resolution timer function `omp_get_wtime()`.

1.3 Results

1.3.1 Output

Result below shows the 20 cluster centroids to which the algorithm converged

```
In [100]: kmeans_op = pd.read_csv('../openmp/result.csv')
          kmeans_op.head(20)
```

```
Out[100]:
```

	CLUSTER_ID	NUM_PTS	CENTROID_X	CENTROID_Y
0	0	25000	7507.21508	7591.16904
1	1	25000	20986.22856	7583.28724
2	2	25000	34505.74344	7590.43368
3	3	25000	47995.84272	7603.73724
4	4	25000	61505.01472	7615.65620
5	5	25000	-7509.70652	7600.08716
6	6	25000	-20999.28700	7604.38880
7	7	25000	-34503.96696	7606.82912
8	8	25000	-47997.90336	7590.85656
9	9	25000	-61505.04916	7603.93612
10	10	25000	-7496.48468	-7599.07976
11	11	25000	-21017.49684	-7574.37808
12	12	25000	-34508.74520	-7605.09100
13	13	25000	-47999.49952	-7606.13836

14	14	25000	-61518.25144	-7600.48044
15	15	25000	7505.19200	-7591.39480
16	16	25000	21011.26088	-7583.93400
17	17	25000	34495.05208	-7616.55264
18	18	25000	48001.69036	-7579.16096
19	19	25000	61501.02676	-7603.22488

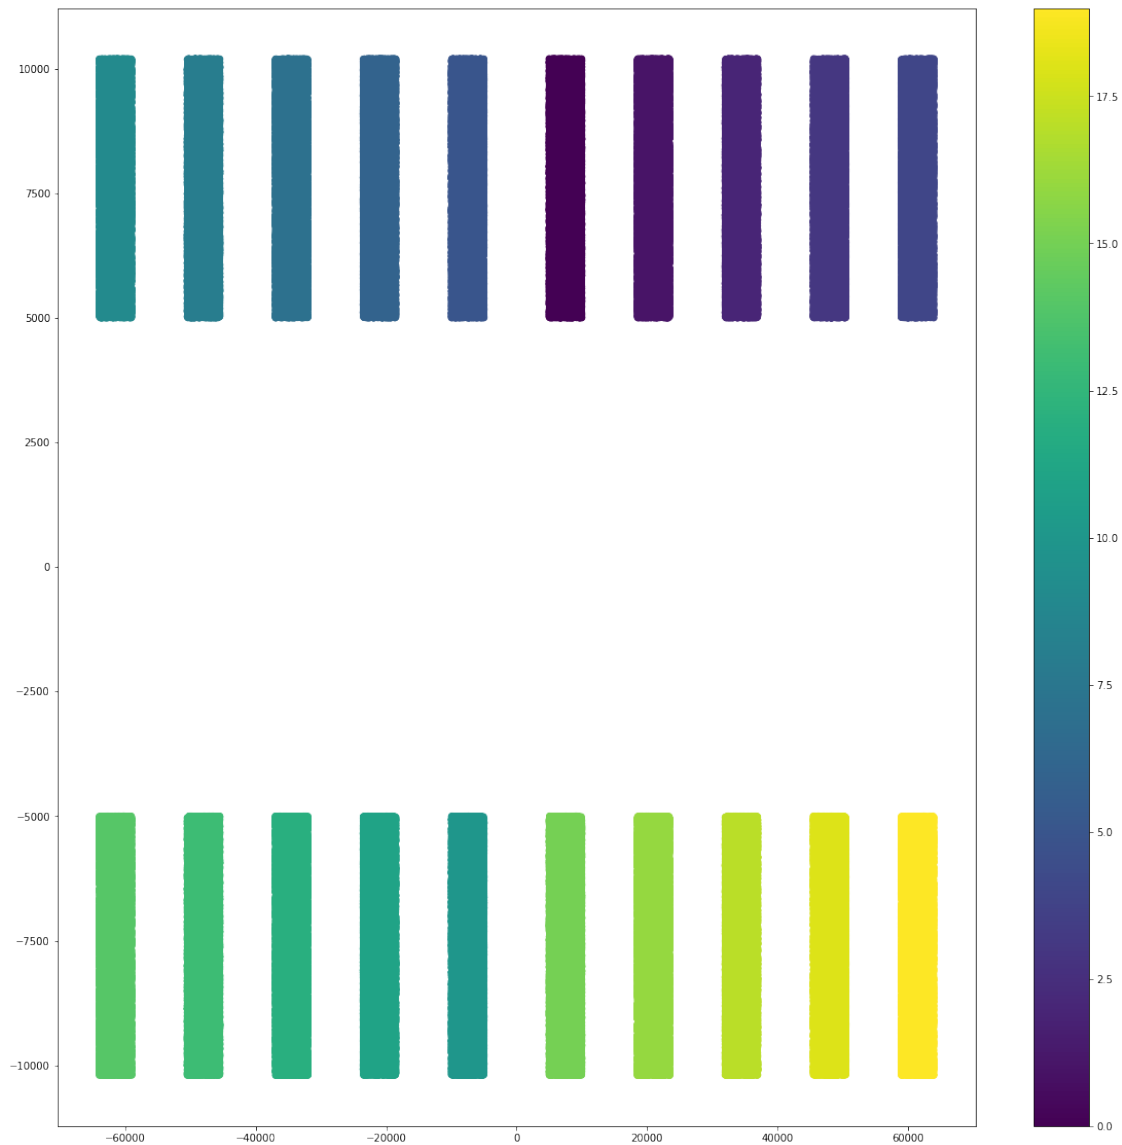
Figure below shows the different clusters formed in different clusters after the program is run.

```
In [101]: clusters_formed = pd.read_csv("../openmp/clusters_formed.csv", header = None)
```

1.3.2 Visualization of the clusters formed

```
In [21]: plt.figure(figsize = (20,20))
plt.scatter(data[0], data[1], c = clusters_formed[0])
plt.colorbar()
```

```
Out[21]: <matplotlib.colorbar.Colorbar at 0x7f3da4618710>
```



1.3.3 Execution times in msecs

```
In [3]: openmp_timings = pd.read_csv('../openmp/openmp_timings.csv')
        openmp_timings.head()
```


```
Out[3]:
```

	seq	static 4	dynamic 4	static 8	dynamic 8	static 16	dynamic 16
0	537.630	191.073	292.110	175.893	249.565	135.812	248.685
1	527.383	184.151	250.661	196.138	253.082	123.065	247.491
2	568.125	193.157	244.271	171.515	291.176	131.803	248.790
3	528.254	263.305	291.080	175.525	253.013	130.494	246.350
4	564.453	225.322	294.039	132.152	252.753	149.175	250.485

1.3.4 Plots

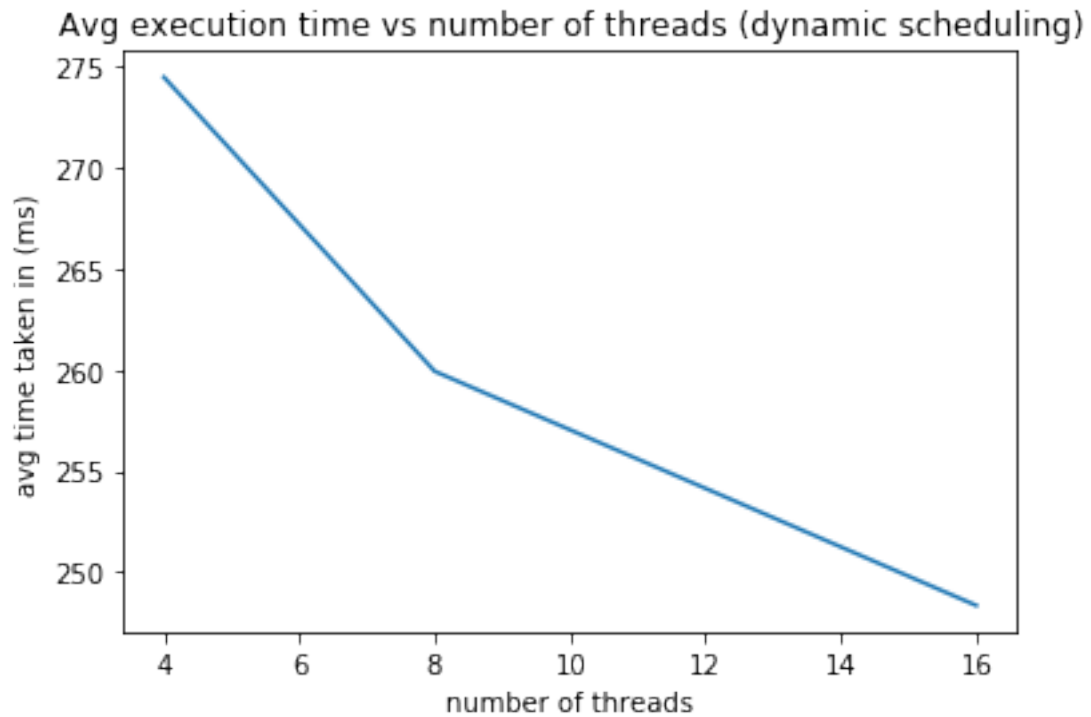
```
In [4]: openmp_avg_times = openmp_timings.mean().values
        seq_time = openmp_avg_times[0]
        openmp_avg_times = openmp_avg_times[1:]
        dynamic_idx = [1,3,5]
        static_idx = [0,2,4]
        nthreads = [4,8,16]
        plt.plot(nthreads, openmp_avg_times[static_idx])
        plt.xlabel('number of threads')
        plt.ylabel('avg time taken in (ms)')
        plt.title('Avg execution time vs number of threads (static scheduling)')
```

```
Out[4]: Text(0.5,1,'Avg execution time vs number of threads (static scheduling)')
```



```
In [5]: plt.plot(nthreads, openmp_avg_times[dynamic_idx])
        plt.xlabel('number of threads')
        plt.ylabel('avg time taken in (ms)')
        plt.title('Avg execution time vs number of threads (dynamic scheduling)')
```

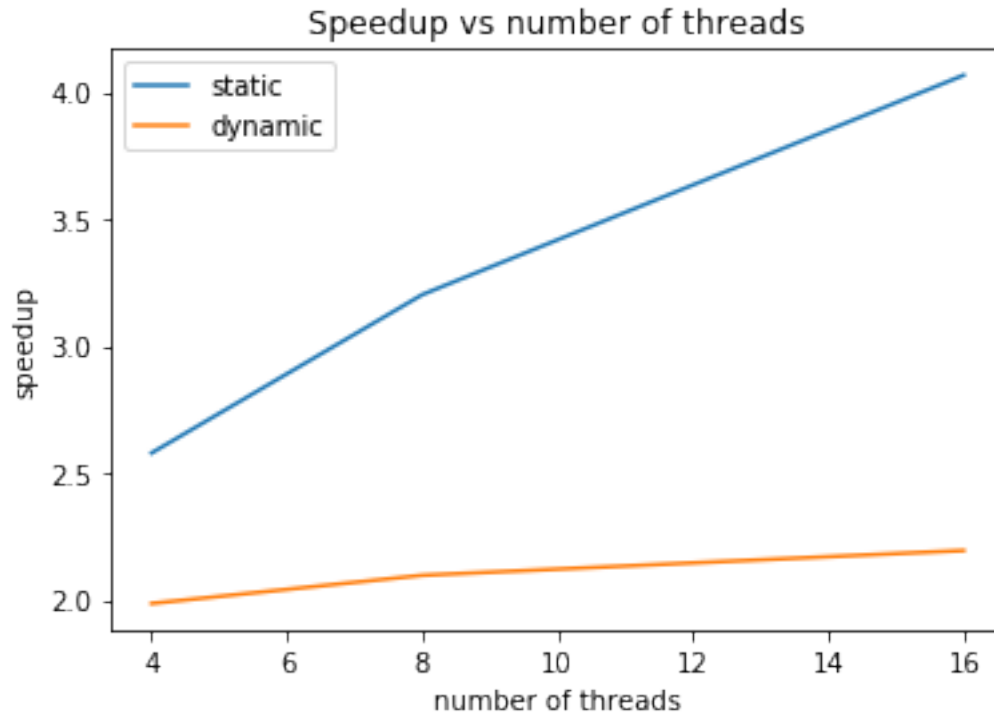
```
Out[5]: Text(0.5,1,'Avg execution time vs number of threads (dynamic scheduling)')
```



```
In [6]: speedup_static = [seq_time/openmp_avg_times[i] for i in static_idx]
        speedup_dynamic = [seq_time/openmp_avg_times[i] for i in dynamic_idx]
```

```
plt.figure()
plt.plot(nthreads, speedup_static)
plt.plot(nthreads, speedup_dynamic)
plt.xlabel('number of threads')
plt.ylabel('speedup')
plt.title('Speedup vs number of threads')
plt.legend(['static', 'dynamic'])
```

```
Out[6]: <matplotlib.legend.Legend at 0x7f3ba7127f28>
```



1.4 Observations

1. Due to well structured data with proper choice of initial centroids we obtain the desired clusters in less number of iterations.
2. Every iteration of sequential Kmeans is roughly $O(np_{\text{oints}} * n_{\text{clusters}})$, and overall running sequential Kmeans on 500K points takes on an avg 520 msec.
3. Parallelizing the loop with openmp can give substantial speedups for eg. for 4 threads we obtain almost 2.5x speedup
4. With increasing number of threads we obtain higher speedups however after a certain point we can expect to hit a saturation point after which the curve saturates. This is in accordance to Amdahl's law.
5. The speedups obtained through dynamic scheduling are lesser as compared to static, this is because the workload(parallelized loop) is regular among the threads hence we dont see much performance boost.

2 MPI Merge Sort

2.1 Methodology

In this we implement a parallel merge sort program in mpi that can sort a list of integers. The idea here is process 0 first gets the data in global array this data is then scattered to the individual

processes local array using mpi function `MPI_Scatterv()`, the reason I have used `MPI_Scatterv` is because here the data may not be divided equally among the processors in case the number of integers in the array is not divisible by the number of processes. In such a case each processor first gets $N/nprocs$ number of integers and any leftover integers are assigned equally among the processors instead of assigning all remaining ints to one processor. Then the `parallel_merge_sort()` function is called for each of the threads. The parallel merge sort works as follows:

1. Each of the processes first sort there part of the local array using `stl` function quick sort.
2. After that a loop runs for $\log_2(nprocs)$ number of levels
3. Each of the processor finds if it is the root/parent processor if it is not then it sends its data (local array and size) to the root processor using an `MPI_Send`.
4. The root processor receives the data from its partner using `MPI_Recv` and merges the two arrays and updates its array to the merged array and also updates its size.
5. Finally the result is obtained in processor 0.

2.2 Experimental Setup

Input: The input data file `merge_data.txt` contains 1L integers. Output: The output of the algorithm is sorted list of integers along with the execution time for running the program The output is generated in `result.csv` file

The experiment was conducted for 8, 16, 32, 64 and 80 processes and timings for each were noted and compared with sequential quick sort. The total execution time for the algorithm is obtained by calculating the maximum time across all the processes.

2.3 Results

2.3.1 Output

For the sorted output please refer the `result.csv` file

2.3.2 Execution times in msec

```
In [7]: sort_times = pd.read_csv('../mpi/mpi_result.csv')
        sort_times.head()
```

```
Out[7]:
```


	seq sort	8 procs	16 procs	32 procs	64 procs	80 procs
0	15.173	5.115271	8.249760	13.636112	27.039051	6.471157
1	15.606	5.060673	8.991241	9.444475	11.289597	16.919613
2	15.156	5.105257	8.766890	14.188051	11.272907	15.195131
3	15.062	5.049944	13.677359	15.269756	18.224955	14.958858
4	14.891	5.011320	14.421225	19.699812	12.118101	27.476549

2.3.3 Plots

```
In [8]: mpi_avg_times = sort_times.mean().values
        seq_time = mpi_avg_times[0]
        mpi_avg_times = mpi_avg_times[1:]
        nprocs = [8,16,32,64,80]
        plt.plot(nprocs, mpi_avg_times)
```


```
plt.xlabel('number of processes')
plt.ylabel('avg time taken in (ms)')
plt.title('Avg execution time vs number of processes')
```

Out[8]: Text(0.5,1,'Avg execution time vs number of processes')



```
In [9]: speedup = [seq_time/i for i in mpi_avg_times]
plt.plot(nprocs, speedup)
plt.xlabel('number of processes')
plt.ylabel('Speedup')
plt.grid()
plt.title('Speedup vs number of processes')
```

Out[9]: Text(0.5,1,'Speedup vs number of processes')



2.4 Observations

1. C library quicksort is highly optimized and has a linearithmic running time and takes on an average 15 msecs to run.
2. For the given 100K dataset maximum speedup is obtained with 8 processes, with a speedup factor of almost 3.
3. With increasing number of processes the communication overhead overshadows any performance benefits that we may expect to see, therefore the speedup curve decreases with increasing number of processes.

3 Reduction using CUDA

3.1 Methodology

Here we implement tree based reduction of a array of integers using CUDA. Two different versions of reduction is implemented one using global memory and one using shared memory. As here we are working with 20K ints and using only one thread block and 1K threads, each thread first reduces 20 elements each, after all the threads are done we are left with 1K ints on which we apply the conventional tree based reduction algorithm to get the result. The program goes as follows:

1. Allocate memory and copy data to device memory.
2. Launch the kernel with 1 block and 1K threads and arguments as d_data(device data array), res (result array), nprocs (number of processes).

3. Each thread reduces 20 elems each at the end of which we synchronize.
4. Then we start with $N/2$ threads and reduce two elements at a time until we are left with a single thread.
5. Finally the result is stored in the `res[0]`.
6. Copy back the result from device to host.

3.2 Experimental Setup

Input: The input data file `reduce_data.txt` contains 20K ints. Output: The output of the program is a single integer containing the reduction result along with the execution time.

The experiment was conducted twice once for the `res` array allocated in global memory and once for the `res` array allocated in shared memory. The execution times obtained in two cases were recorded and compared.

3.3 Result

3.3.1 Output

Cuda Reduction result: 637144 Actual answer: 637144

3.3.2 Execution times in msec

```
In [112]: reduction_times = pd.read_csv('../cuda/cuda_results_latest.csv')
         reduction_times.head() #time in msec
```

```
Out[112]:
```

	seq	w.o shared	shared
0	0.164458	0.156192	0.146400
1	0.150597	0.153728	0.147264
2	0.150415	0.158784	0.151904
3	0.150508	0.154688	0.150112
4	0.150488	0.151392	0.148224


```
In [113]: reduction_times.mean()
```

```
Out[113]: seq          0.153293
         w.o shared    0.154957
         shared        0.148781
         dtype: float64
```

3.3.3 Plots

```
In [114]: y = reduction_times.mean().values
         x = ['sequential reduction', 'w.o. shared memory', 'with shared memory']
         plt.bar(x,y)
         plt.ylabel('Time in msec')
         plt.title('Comparison of avg execution time')
```


```
Out[114]: Text(0.5,1,'Comparison of avg execution time')
```



```
In [115]: speedup = [y[0]/i for i in y[1:]]
          print(speedup)
          x = ['w.o. shared memory', 'with shared memory']
          plt.bar(x, speedup)
          plt.ylabel('Speedup')
          plt.grid()
          plt.title('Comparison of speedup with and without shared memory')
```

```
[0.9892641045762435, 1.0303291822600769]
```

```
Out[115]: Text(0.5,1,'Comparison of speedup with and without shared memory')
```



3.4 Observation

1. Since the dataset size is only 20K and parallelization has been done with 1000 threads in 1 block the performance of sequential reduction and parallel reduction using cuda is almost similar.
2. We can further optimize the performance of CUDA program by using shared memory, using shared memory the speedup improves from 0.9892 to 1.03